

# 10: Additional Workloads

## 10.1: Manual Run Jobs

**Jobs** are a handy way to execute “run to completion” style workloads, unlike Deployments which are meant to run forever until they are terminated by an error or by a user. Let’s explore running the job below which calculates the value of pi:

lab10-job.yaml

```
1  apiVersion: batch/v1
2  kind: Job
3  metadata:
4    name: pi
5  spec:
6    template:
7      metadata:
8        name: pi
9      spec:
10       securityContext:
11         runAsUser: 1000
12       containers:
13         - name: pi
14           image: perl:5.34.0
15           command: ["perl", "-Mbignum=bpi", "-wle", "print
16             bpi(2000)"]
17       restartPolicy: Never
```

Download and apply lab10-job.yaml.

```
curl <url_of_yaml_above> -o lab10-job.yaml
kubectl apply -f lab10-job.yaml
```

Notice the following observations:

- After the container exits successfully Kubernetes does not try to start another container like it does with Deployments

- The Pod and Job stick around after completion so that you can view the output
- Find the output calculated by the pi job using kubectl commands (get, describe, logs)

Try applying the yaml again with kubectl, does the job run a second time? What would you need to do if you wanted to rerun the job?

## 10.2: Parallel Jobs with a Work Queue

Parallel jobs with a work queue can create several pods which coordinate with themselves or with some external service which part of the job to work on.

If your application has a work queue implementation for some remote data storage, for example, this type of Job can create several parallel worker pods that will independently access the work queue and process it.

Parallel jobs with a work queue come with the following features and requirements:

- for this type of Job, you should leave `.spec.completions` unset.
- each worker pod created by the Job is capable of assessing whether or not all its peers are done and, thus, the entire Job is done (e.g. each pod can check if the work queue is empty and exit if so).
- when any pod terminates with success, no new pods are created.
- once at least one pod has exited with success and all pods are terminated, then the job completes with success as well.
- once any pod has exited with success, other pods should not be doing any work and should also start exiting.

Let's add parallelism to see how these types of Jobs work, notice the `.spec.parallelism` field was added and set to 3.

## lab10-parajob.yaml

```
1  apiVersion: batch/v1
   kind: Job
2  metadata:
     name: primes-parallel-wq
3  labels:
     app: primes
4  spec:
     parallelism: 3
5  template:
     metadata:
6     name: primes
     labels:
7     app: primes
     spec:
8     securityContext:
         runAsUser: 1000
9     containers:
1     - name: primes
0       image: debian:stable-slim
1       command: ["bash"]
1       args:
1       [
2         "-c",
1         "current=0; max=110; echo 1; echo 2; for((i=3;i<=max;)); do
3 for((j=i-1;j>=2;)); do if [ `expr $i % $j` -ne 0 ] ; then current=1; else
1 current=0; break; fi; j=`expr $j - 1`; done; if [ $current -eq 1 ] ; then echo $i;
4 fi; i=`expr $i + 1`; done",
1       ]
5     restartPolicy: Never
1
6
1
7
1
8
1
9
2
2
0
2
1
2
2
2
2
3
2
4
2
5
```

2  
6

Now, let's open two terminal windows. In the first terminal, watch the pods:

```
kubect1 get pods -l app=primes -w
```

Download and apply lab10-job.yaml using a second terminal.

```
curl <url_of_yaml_above> -o lab10-parajob.yaml  
kubect1 apply -f lab10-parajob.yaml
```

Next, let's see what's happening in the first terminal window:

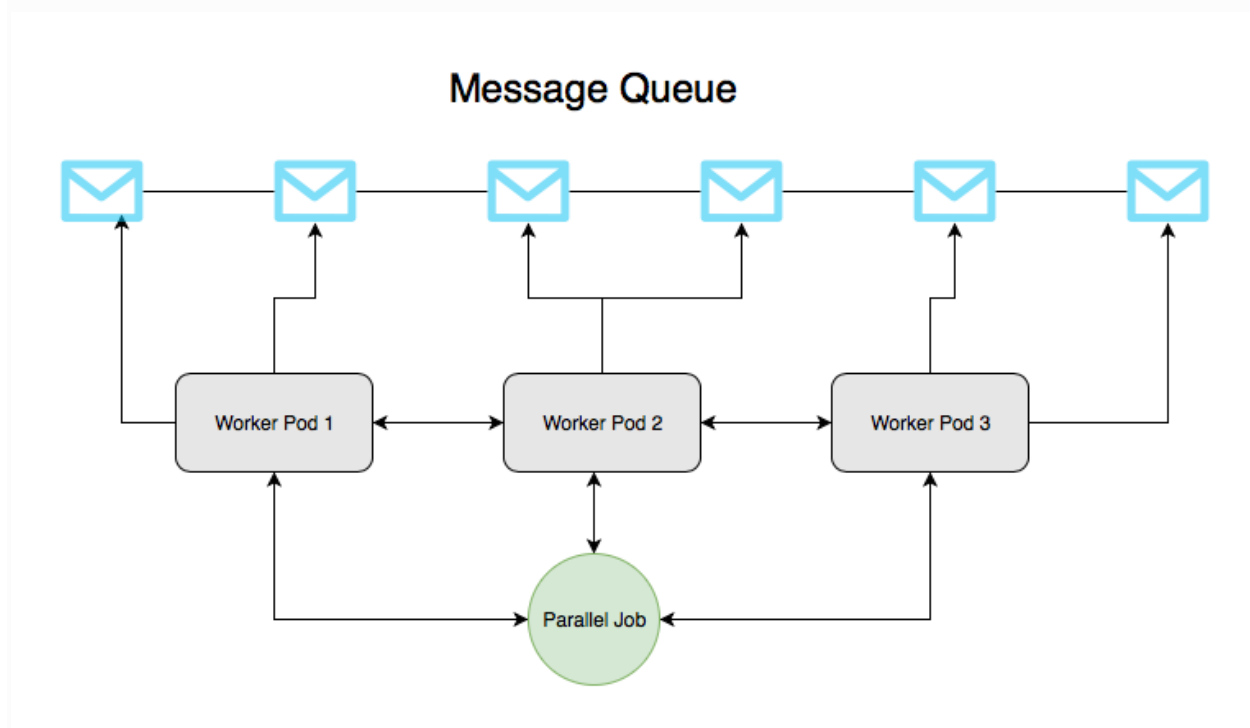
```
kubect1 get pods -l app=primes -w
```

```
*****SAMPLE OUTPUT*****  
NAME                READY    STATUS    RESTARTS   AGE  
primes-parallel-b2whq 0/1      Pending  0           0s  
primes-parallel-b2whq 0/1      Pending  0           0s  
primes-parallel-vhvqm 0/1      Pending  0           0s  
primes-parallel-cdfdx 0/1      Pending  0           0s  
primes-parallel-vhvqm 0/1      Pending  0           0s  
primes-parallel-cdfdx 0/1      Pending  0           0s  
primes-parallel-b2whq 0/1      ContainerCreating 0           0s  
primes-parallel-vhvqm 0/1      ContainerCreating 0           0s  
primes-parallel-cdfdx 0/1      ContainerCreating 0           0s  
primes-parallel-b2whq 1/1      Running   0           4s  
primes-parallel-cdfdx 1/1      Running   0           7s  
primes-parallel-vhvqm 1/1      Running   0          10s  
primes-parallel-b2whq 0/1      Completed  0          17s  
primes-parallel-cdfdx 0/1      Completed  0          21s  
primes-parallel-vhvqm 0/1      Completed  0          23s
```

As you see, the kubect1 created three pods simultaneously. Each pod was calculating the prime numbers in parallel and once each of them completed the task, the Job was successfully completed as well.

In a real-world scenario, we could imagine a Redis list with some work items (e.g messages, emails) in it and three parallel worker pods created by the Job (see the Image above). Each pod could have a script to requests a new message from the list, process it, and check if there are more work items left.

If no more work items exist in the list, the pod accessing it would exit with success telling the controller that the work was successfully done. This notification would cause other pods to exit as well and the entire job to complete. Given this functionality, parallel jobs with a work queue are extremely powerful in processing large volumes of data with multiple workers doing their tasks in parallel.



Example: <https://kubernetes.io/docs/tasks/job/fine-parallel-processing-work-queue/>

## 10.3: Scheduled Jobs¶

In the above scenario we ran a job manually and once. Sometimes you want to run a job at a regular scheduled interval or once at a particular time. [CronJobs](#) in Kubernetes provide this functionality.

## lab10-cronjob.yaml

```
1  apiVersion: batch/v1
2  kind: CronJob
3  metadata:
4    name: hello
5  spec:
6    schedule: "*/1 * * * *"
7    jobTemplate:
8      spec:
9        template:
10       spec:
11         securityContext:
12           runAsUser: 1000
13         containers:
14           - name: hello
15             image: busybox
16             args:
17               - /bin/sh
18               - -c
19               - date; echo Hello from the Kubernetes
20 cluster
      restartPolicy: OnFailure
```

Download and apply lab10-cronjob.yaml.

```
curl <url_of_yaml_above> -o lab10-cronjob.yaml
```

```
kubectl apply -f lab10-cronjob.yaml
```

This will schedule the job to run every minute. We watch this in action by both view the job or watching the job:

\*\*\* EXAMPLE OUTPUT \*\*\*

```
$ kubectl get cronjob hello
```

NAME	SCHEDULE	SUSPEND	ACTIVE	LAST SCHEDULE	AGE
hello	*/1 * * * *	False	0	<none>	9s

```
$ kubectl get jobs --watch
```

NAME	DESIRED	SUCCESSFUL	AGE
hello-1543405740	1	1	23s

